



DATA PERSISTENCE

NEUEDA.COM

Objectives

- Data Persistence in OpenShift
- Persistent Volumes using NFS
- Persistent Volume Claims
- OpenShift Volumes
- Data Persistence Using Databases
- OpenShift Databases

Data Persistence in OpenShift

- In our previous lab we were able to deploy an application, but any data we put in was lost when the app updated.
- Actually, by design Kubernetes Pods are *ephemeral*.
 - Any given Pod could be replaced at any time.
- However persistent state is a critical aspect of many applications.

Example: Pods are Ephemeral

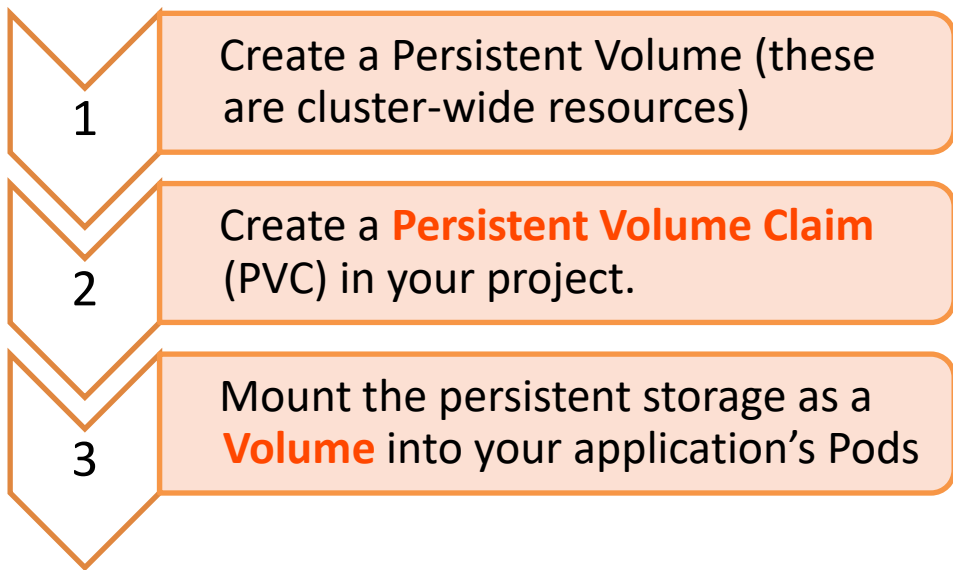
- Let's look at the Image Uploader deployment again
- Upload an image and see how it displays in the website
- Now scale the application down to 0 pods, then scale back up to 1 pod
- Open the URL again - no image!
- What happened? Pods (and their memory) are ephemeral
 - When the Pod is deleted, all its memory goes with it.

Handling Permanent Data

- OpenShift makes persistent storage available for
 - Data that is shared between Pods
 - Data that needs to persist past the lifetime of any particular Pod
- Permanent storage in pods is handled using **Persistent Volumes** (PVs)
- This can be supplied from a variety of storage solutions
 - NFS, HostPath (local directories), AWS EBS, Azure Disk, etc

Persistent Volumes using NFS

- In OpenShift the process of creating persistent file storage has three steps:



Creating a Persistent Volume

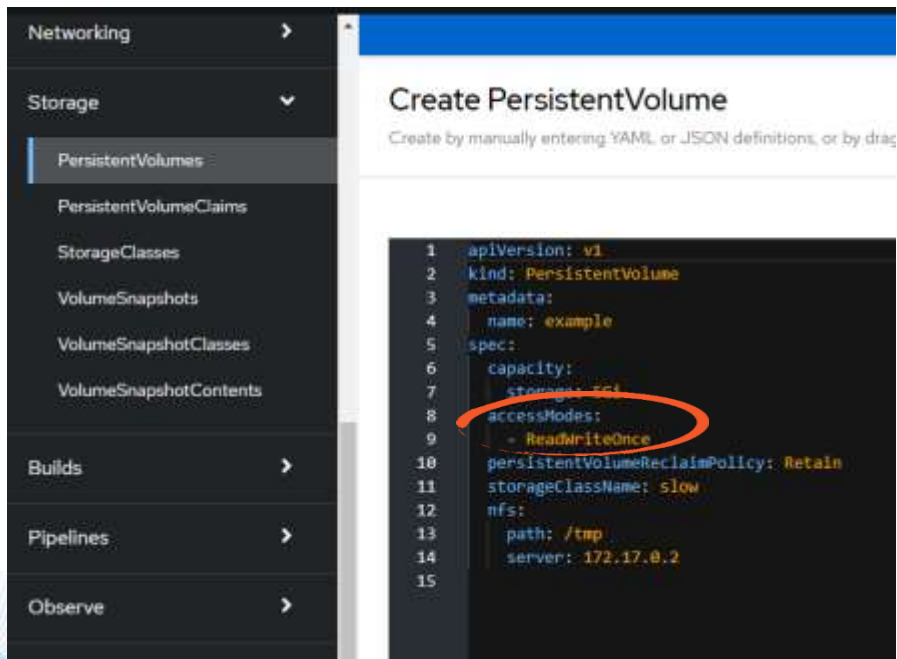
- Persistent Volumes are cluster-wide resources and really should only be allocated by the Cluster Administrator
- Go to the Administrator perspective in your Web Console and click on the Storage item.
 - See there is no “PersistentVolumes” option.
- In the Cluster Admin console you can see that option is available.

Creating a Persistent Volume

The screenshot shows the Red Hat OpenShift Container Platform web console. The left sidebar has a 'Storage' section with 'PersistentVolumes' highlighted. The main content area is titled 'PersistentVolumes' and features a 'Create PersistentVolume' button in the top right corner. Below the title is a search bar and a table of existing Persistent Volumes.

Name	Status	Claim	Capacity	Labels	Created
pvc-5b09e0dc-23c7-4dc6-a820-	Bound	pvc-91d2ae1d9d	1Gi	topology.k...=eu... topology.k...=eu...	4 May 2022, 09:45

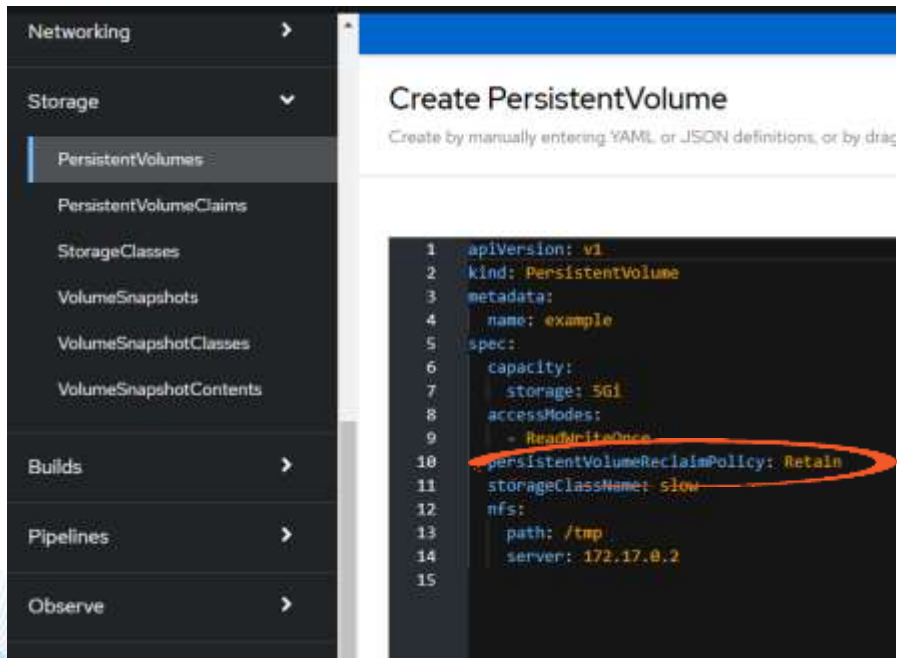
Configuring a PersistentVolume: accessMode



```
1  apiVersion: v1
2  kind: PersistentVolume
3  metadata:
4    name: example
5  spec:
6    capacity:
7      storage: 1Gi
8    accessModes:
9      - ReadWriteOnce
10   persistentVolumeReclaimPolicy: Retain
11   storageClassName: slow
12   nfs:
13     path: /tmp
14     server: 172.17.0.2
```

- The **accessMode** can be
 - Read/Write once (RWO) - Can be mounted as read/write by a single node in the cluster
 - Read-only many (ROX) - Can be mounted as read-only by many nodes
 - Read/Write many (RWX) - Can be mounted as read/write by multiple nodes in the cluster.

Configuring a PersistentVolume: Reclaim Policy



Networking >

Storage >

PersistentVolumes

PersistentVolumeClaims

StorageClasses

VolumeSnapshots

VolumeSnapshotClasses

VolumeSnapshotContents

Builds >

Pipelines >

Observe >

Create PersistentVolume

Create by manually entering YAML, or JSON definitions, or by drag

```
1  apiVersion: v1
2  kind: PersistentVolume
3  metadata:
4    name: example
5  spec:
6    capacity:
7      storage: 5Gi
8    accessModes:
9      - ReadWriteOnce
10   persistentVolumeReclaimPolicy: Retain
11   storageClassName: slow
12   nfs:
13     path: /tmp
14     server: 172.17.0.2
15
```

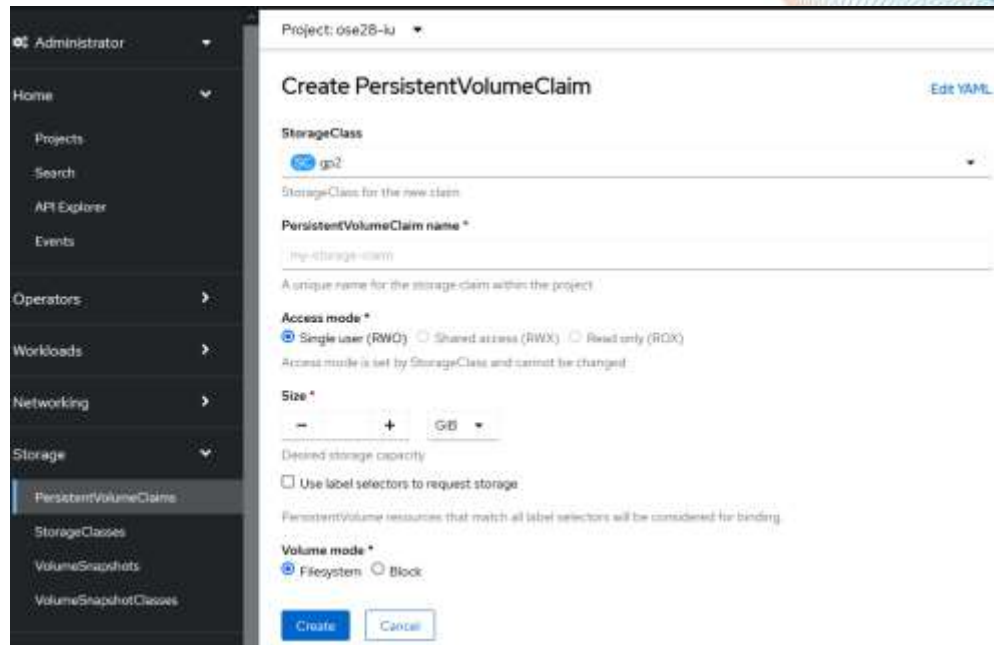
- The reclaim policy dictates how a PV handles reclaiming space after a storage claim on the PV is no longer required. It can be
 - Retain - all data is retaining in the PV, you have to reclaim space manually
 - Recycle - all data is automatically removed when the claim is deleted.

Using Persistent Storage

- To take advantage of a PV, you need to make a claim on that PV.
- PVs represent available storage, and PVCs represent an application's need for that storage.
- When you create a PVC, OpenShift looks for the best fit among the available PVs and reserves it for use by the PVC.
 - PV size vs PVC need - use the smallest available PV
 - Access Mode - use a PV with the same or greater access privileges than that required by the PVC

Adding a Persistent Volume Claim

- PersistentVolumeClaims can be created by accessing the PersistentVolumeClaims item under storage in the Administrator perspective



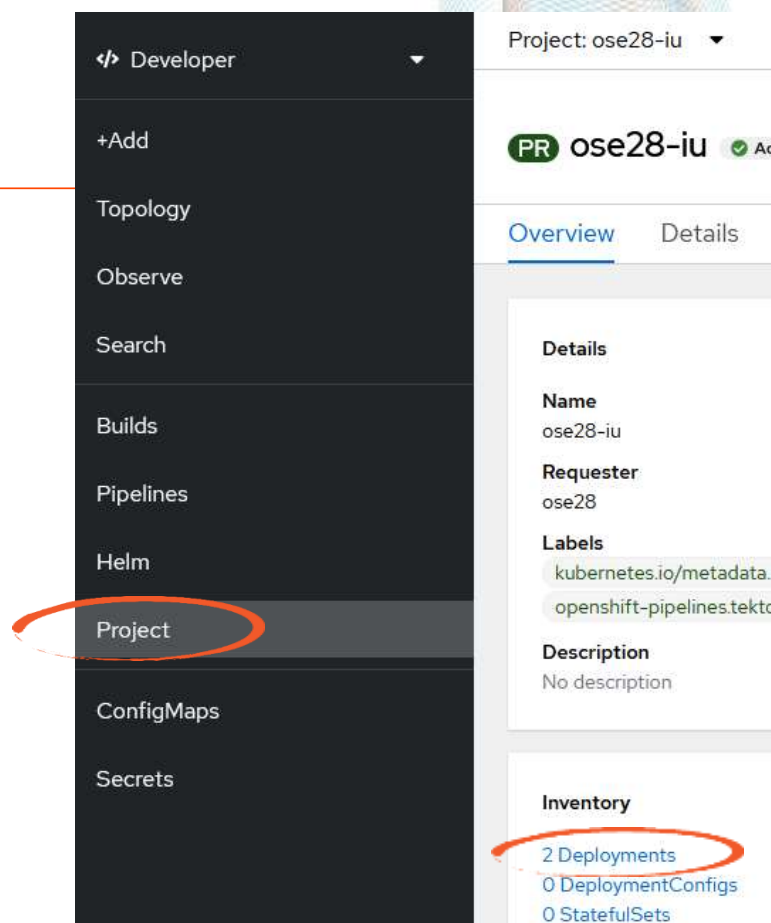
The screenshot displays the 'Create PersistentVolumeClaim' form in the Kubernetes Administrator UI. The left sidebar shows the navigation menu with 'Storage' expanded and 'PersistentVolumeClaims' selected. The main form area is titled 'Create PersistentVolumeClaim' and includes the following fields and options:

- Project:** ose28-lu
- StorageClass:** gp2 (with a dropdown arrow)
- StorageClass for the new claim:** (empty field)
- PersistentVolumeClaim name ***: my-storage-claim (with a text input field)
- Access mode ***: ☒ Single user (RWX) ☐ Shared access (RWX) ☐ Read only (ROX). A note below states: 'Access mode is set by StorageClass and cannot be changed.'
- Size ***: A numeric input field with a '+' button and a 'GB' unit dropdown.
- Desired storage capacity:** ☐ Use label selectors to request storage.
- Volume mode ***: ☒ Filesystem ☐ Block.

At the bottom of the form are 'Create' and 'Cancel' buttons. A note at the bottom of the form states: 'PersistentVolume resources that match all label selectors will be considered for binding.'

OpenShift Volumes

- A *volume* is any filesystem, file or data mounted into an application's Pods to provide persistent data.
- To attach a PVC as a volume in our app using the Web Console you need to get to the project overview page and select the Deployment.



Creating a Volume

The screenshot displays the Red Hat OpenShift Container Platform interface. On the left is a dark sidebar with navigation options: Developer, +Add, Topology, Observe, Search, Builds, Pipelines, and Helm. The main header shows the Red Hat logo and 'OpenShift Container Platform' on the left, and a user profile 'ose28' on the right. The main content area is titled 'Project: ose28-iu' and shows 'Deployments > Deployment details' for a deployment named 'image-uploader'. Below this is a tabbed interface with 'Details' selected. The 'Deployment details' section shows a blue circle with '1 Pod' inside. An 'Actions' dropdown menu is open on the right, with 'Add storage' highlighted by a red circle. Other options in the menu include 'Edit Pod count', 'Pause rollouts', 'Add Health Checks', 'Add HorizontalPodAutoscaler', 'Edit update strategy', and 'Edit resource limits'.

Creating a Volume

- Just select your existing PVC
- The Mount Path is the path within each pod at which the persistent volume should be available - it is application dependent.
- On “Save” the volume is attached.

Red Hat OpenShift Container Platform

Project: ose28-lu

Add Storage to [image-uploader](#)

PersistentVolumeClaim *

☒ Use existing claim

[PVC](#) ose28-pvc

☐ Create new claim

Mount path *

/opt/app-root/src/uploads

Mount path for the volume inside the container.

☐ Mount as read-only

Subpath

Optional path within the volume from which it will be mounted into the container. Defaults to the root of the volume.

The volume will be mounted into all containers. You can [select specific containers instead](#).

Save Cancel

Lab 3: Add Persistent Data to the Image Uploader

- This is a quick lab to implement what we have just seen in your own project.
- Go to [/labs/3-image-uploader-persistent-storage.md](#)

Volumes Separate Data

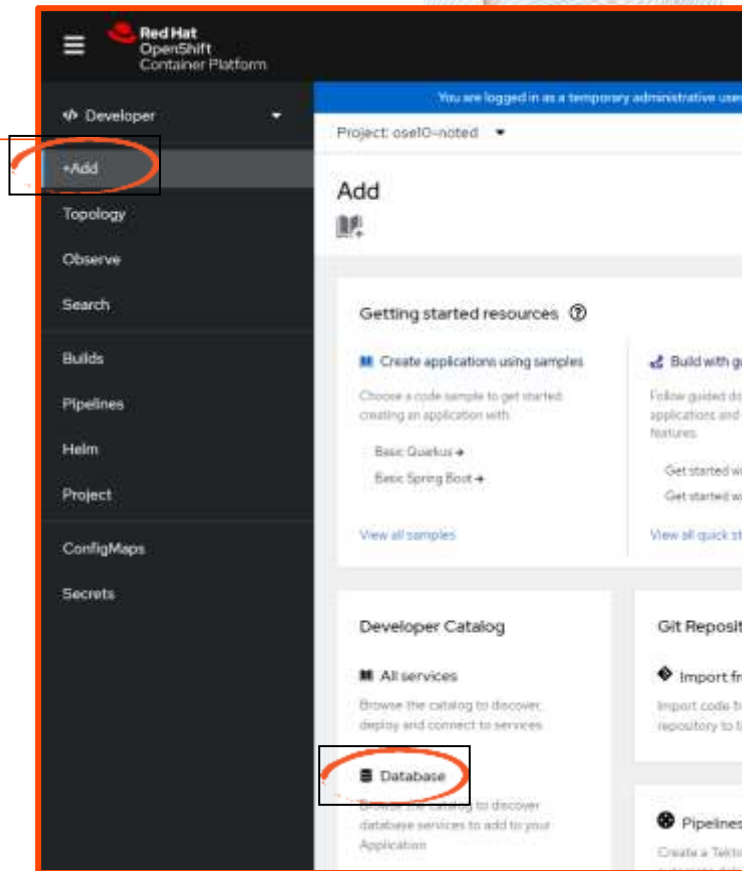
- Suppose everyone set up Volumes on our cluster - why don't you see everyone else's data?
 - Each application deployment uses its own NFS volume to store data.
 - Each NFS volume is then mounted into its own application's mount space.
 - So the data is always separated on the cluster.

Data Persistence Using Databases

- The Persistent Volume technology is only one way to data persistence.
- You can also hook your application to a database
 - If the database is not running on your cluster, then it supplies its own data persistence
 - If the database is running on your cluster, then data persistence is achieved behind the scenes.

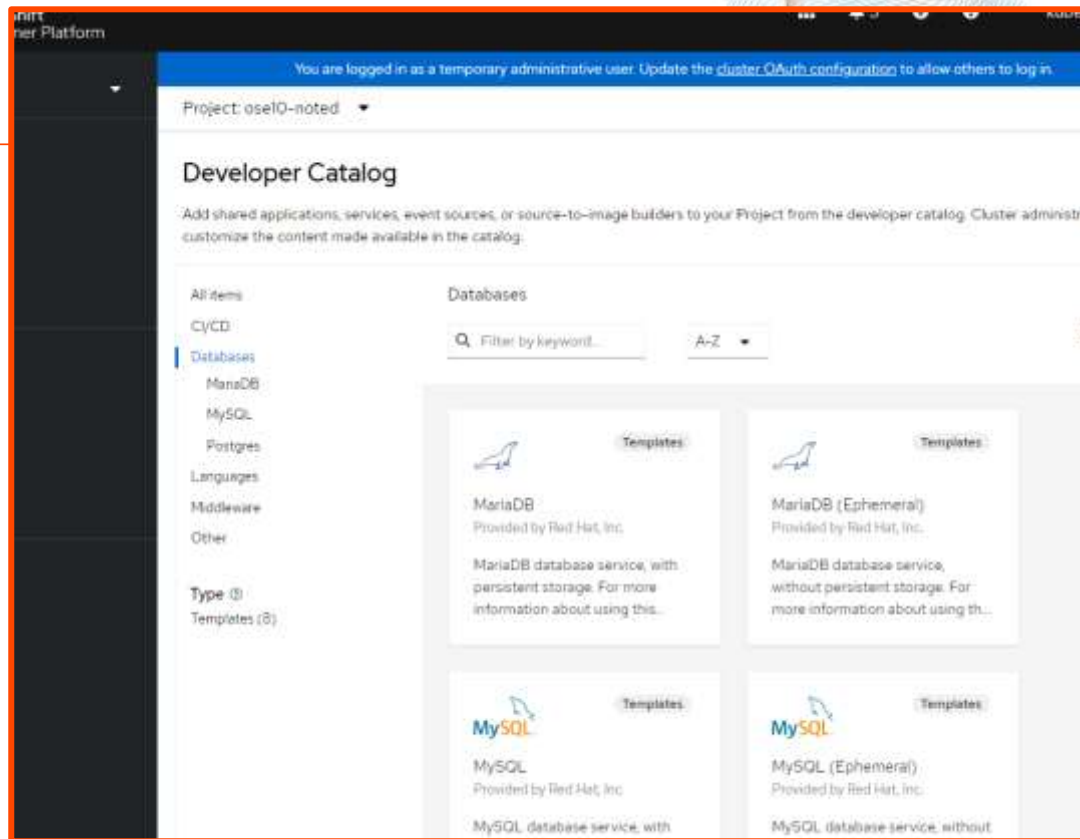
OpenShift Databases

- OpenShift has pre-configured MySQL, MariaDB and PostgreSQL databases available in the Developer Catalogue



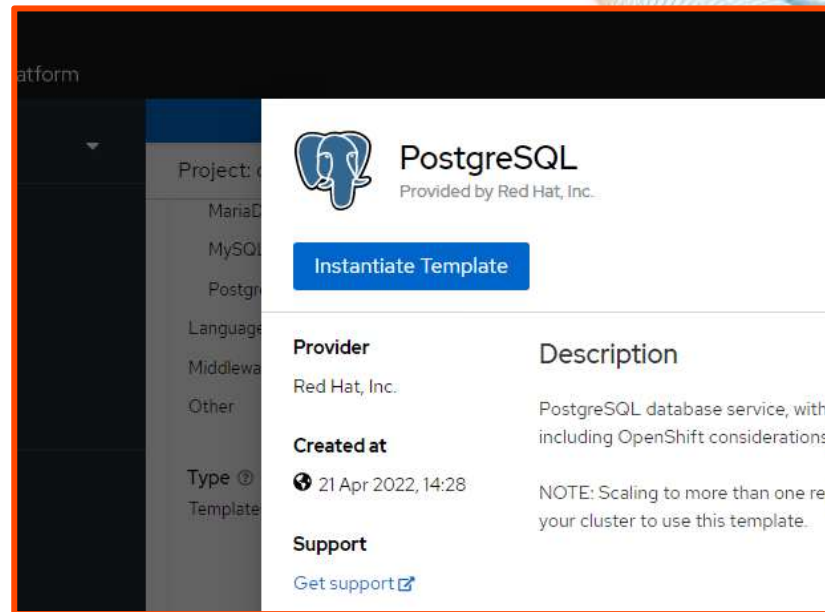
OpenShift Databases

- The databases all offer ephemeral (for testing) and persistent options.



OpenShift Databases

- During Instantiation you can supply a username, password and other configuration details.
- If you use a template you will need to manually configure your app to inject the environment variables needed to connect.



Demo: Petclinic

- In this demo we will deploy a Spring Boot app and a MySQL database on the same cluster and link the two.

Connecting to Databases

- The database deployment we saw is quite manual, OpenShift has a number of more automatic connection methods.
- One of the simplest of these is using a **Secrets** object.
 - Secrets objects store one or more values that are intended to be obscured (like passwords, or certificate files).
 - Secrets objects can be made available to Pods in a namespace without divulging the contents to users.
 - We will see Secrets for connecting a database later on.

Connecting the Databases: Service Binding

- Another method for connecting to databases is through the **Service Binding Operator**
- This is one of the built-in operators your cluster admin can install
- The Service Binding Operator (SBO) allows you to quickly bind an instance of a database to an application deployed on OpenShift
 - No need to distribute secrets, config maps, usernames, passwords.
 - But the database needs to be deployed via an operator or a Helm chart configured to what the SBO needs.

Deploying a DB for the SBO

- There are PostgreSQL Operators in the OpenShift OperatorHub, but once again you would have to configure the how the quarkus-backend connects to it manually
- We will install a *new* Operator repository (or OperatorSource)
 - This provides the OpenShift Operator Lifecycle Manager the means to install the PostgreSQL Operator we need
- This is a cluster administrator level task, but worth seeing in a demo!

Service Binding Operator Usage

- There isn't an interface in the console for creating a Service Binding Object, so you have to supply the YAML configuration directly.
- You can either do this through the command line, or using the “+” in the top right of your console.
- This will then make the secrets in the services specified available directly to the application.

Demo: Notes App with PostgreSQL

- Let's go back to our demo notes app and see how we can add a database in OpenShift.

Summary

- Data Persistence in OpenShift
- Persistent Volumes using NFS
- Persistent Volume Claims
- OpenShift Volumes
- Data Persistence Using Databases
- OpenShift Databases

Questions and Comments?

