

TEMPLATES AND OPERATORS

Objectives

- OpenShift Templates
- Templates in the Web Console and the CLI
- Template File Structure
- Kubernetes Operators
- Operators Structure and Operation

Templates and Operators

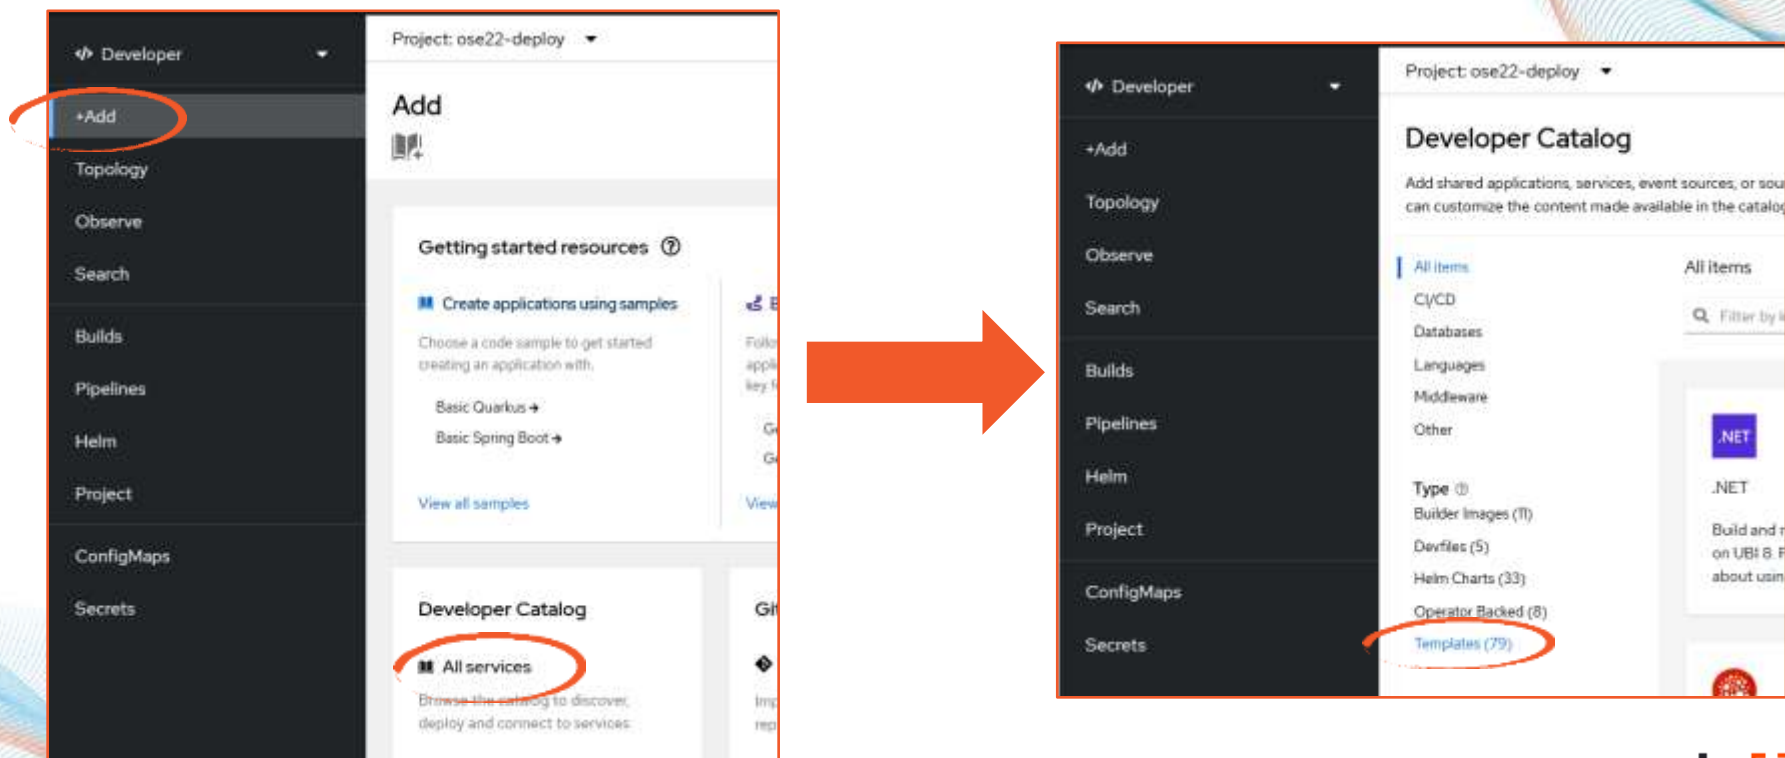
- **Templates** and **Operators** automate repetitive tasks in OpenShift
 - Triggering builds and deployments when source code changes
 - Restarting failed Pods
 - An Operator upgrading your database server
- A Template automates the creation of a set of resources.
- An Operator also deploys an application and its resources then continues to watch and govern those resources.

TEMPLATES

Templates

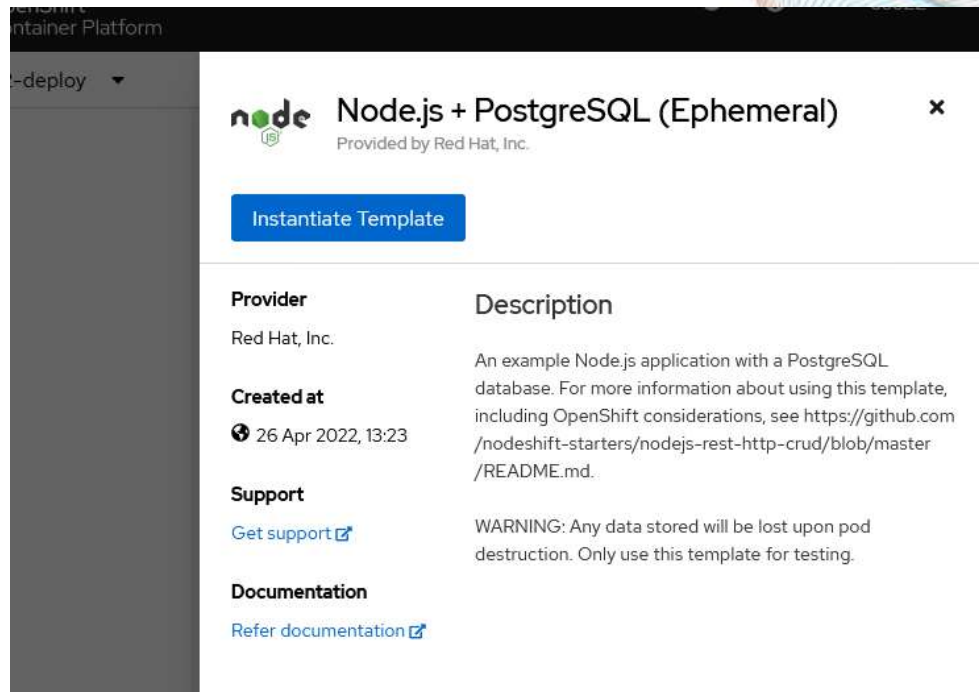
- A template is a list of objects and the names parameters of their configuration
- To view the available templates in the Web Console:
 - Click on the “+Add” item in the Developer Perspective
 - Click on “All services” in the Developer Catalog
 - The click on “Template” in the **Type** submenu in the left hand menu.

Templates in the Web Console



Instantiating a Template in the Web Console

- When a template is selected, a description is shown and you are invited to instantiate it

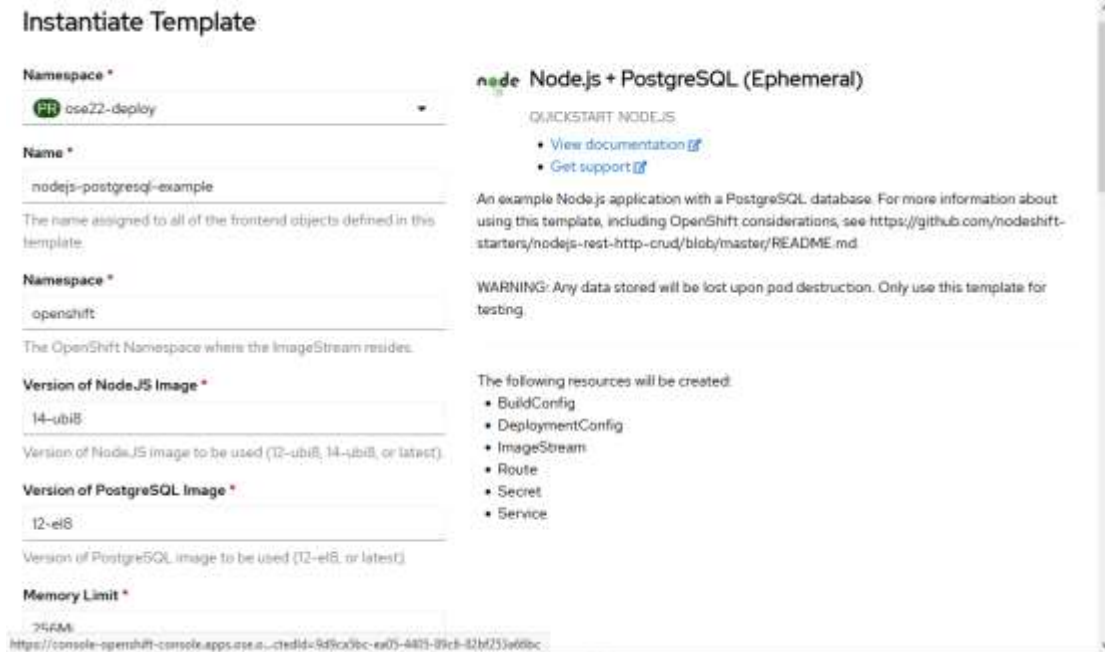


The screenshot displays the OpenShift Container Platform web console interface. At the top, the title bar reads 'OpenShift Container Platform'. Below it, a dropdown menu shows '-deploy'. The main content area features a modal window for the 'Node.js + PostgreSQL (Ephemeral)' template, provided by Red Hat, Inc. The modal has a blue 'Instantiate Template' button. Below the button, there is a table with two columns: 'Provider' and 'Description'. The 'Provider' column lists 'Red Hat, Inc.' and 'Created at' as '26 Apr 2022, 13:23'. The 'Description' column contains a detailed explanation of the template, including a warning about data loss upon pod destruction. Links for 'Support' and 'Documentation' are also visible.

Provider	Description
Red Hat, Inc.	An example Node.js application with a PostgreSQL database. For more information about using this template, including OpenShift considerations, see https://github.com/nodeshift-starters/nodejs-rest-http-crud/blob/master/README.md .
Created at 26 Apr 2022, 13:23	WARNING: Any data stored will be lost upon pod destruction. Only use this template for testing.
Support Get support	
Documentation Refer documentation	

Instantiating a Template in the Web Console

- When you click on “Instantiate Template” you are given a list of configurable parameters:



Instantiate Template

Namespace *
ose22-deploy

Name *
nodejs-postgresql-example
The name assigned to all of the frontend objects defined in this template.

Namespace *
openshift
The OpenShift Namespace where the ImageStream resides.

Version of NodeJS Image *
14-ubi8
Version of NodeJS image to be used (12-ubi8, 14-ubi8, or latest).

Version of PostgreSQL Image *
12-el8
Version of PostgreSQL image to be used (12-el8, or latest).

Memory Limit *
256Mi
https://console-openshift-console.apps.ose.a...ctedId=8dfca9bc-ea05-4475-8fcb-82bf253a66bc

node Node.js + PostgreSQL (Ephemeral)

QUICKSTART NODEJS

- View documentation
- Get support

An example Node.js application with a PostgreSQL database. For more information about using this template, including OpenShift considerations, see <https://github.com/nodehift-starters/nodejs-rest-http-crud/blob/master/README.md>

WARNING: Any data stored will be lost upon pod destruction. Only use this template for testing.

The following resources will be created:

- BuildConfig
- DeploymentConfig
- ImageStream
- Route
- Secret
- Service

Templates in the oc CLI

- You can also access, edit and deploy templates from the command line.
- OpenShift Templates included with a particular cluster install are in the `openshift` namespace. To see the list:

```
$ oc get templates -n openshift
```

NAME	DESCRIPTION	PARAMETERS	OBJECTS
3scale-gateway	3scale's APICast is an NGINX based API gateway used to integrate your interna...	17 (8 blank)	3
amq63-basic	Application template for JBoss A-MQ brokers. These can be deployed as standal...	11 (4 blank)	6
amq63-persistent	An example JBoss A-MQ application. For more information about using this temp...	13 (4 blank)	8
amq63-persistent-ssl	An example JBoss A-MQ application. For more information about using this temp...	18 (6 blank)	12
amq63-ssl	An example JBoss A-MQ application. For more information about using this temp...	16 (6 blank)	10
apicurito	Design beautiful, functional APIs with zero coding, using a visual designer f...	7 (1 blank)	7
cache-service	Red Hat Data Grid is an in-memory, distributed key/value store.	8 (1 blank)	4
cakephp-mysql-example	An example CakePHP application with a MySQL database. For more information ab...	21 (4 blank)	8
>--SNIP--<			

oc process

- The `oc process` subcommand processes a template
 - This produces a valid YAML manifest for the template objects with the specified parameter values filled in.
 - The output is sent to `stdout`, but could equally be piped to `oc create` to actually create the template objects.

```
$ oc process -n openshift nginx-example
```

Viewing the parameters

- The `--parameters` option will display the configurable parameters from the template, in case you need to adjust the defaults
- Parameters can be set in `oc process` using successive `--param` or `-p` arguments

```
$ oc process -n openshift --parameters nginx-example
```

NAME	DESCRIPTION	VALUE
NAME	The name assigned to all of the frontend objects.	nginx-example
NAMESPACE	The OpenShift Namespace where the ImageStream resides.	openshift
NGINX_VERSION	Version of NGINX image to be used (1.20-el8 by default).	1.20-el8
MEMORY_LIMIT	Maximum amount of memory the container can use.	512Mi
SOURCE_REPOSITORY_URL	The URL of the it repository	https://github.com/sclorg/nginx-ex.git
...		

```
$ oc process -n openshift nginx-example -p NAME=nginx-two -p MEMORY_LIMIT=256Mi
```

Instantiating a Template in the CLI

- To instantiate in the CLI
 - Process a template
 - Feed the resulting YAML into the `oc create -f` subcommand

```
$ oc process -n openshift nginx-example -p NAME=nginx-two | oc create -f -  
service/nginx-two created  
route.route.openshift.io/nginx-two created  
imagestream.image.openshift.io/nginx-two created  
buildconfig.build.openshift.io/nginx-two created  
deploymentconfig.apps.openshift.io/nginx-two created  
$ oc get dc nginx-two  
NAME          REVISION  DESIRED  CURRENT  TRIGGERED BY  
nginx-two     0         1        0        config,image/nginx-two:latest
```

This means read from stdin

A Template file is JSON/YAML with these fields:

Metadata

- This section explains what the template does, and adds search tags. It also can select an icon from a list and gives urls for support etc.

Labels

- Labels that are added to each object created when the template is instantiated

objects

parameters

Writing your own OpenShift Template

```
apiVersion: template.openshift.io/v1
kind: Template
```

```
metadata:      # This section explains what the template does, and adds search
                # tags. It also can select an icon from a list and gives links
                # for support etc.
```

```
message:      # A message returned to the user on instantiation
```

```
labels:      # Labels to be added to each object created on instantiation
```

```
objects:      # This is the heart of the template - the list of objects that
                # will be created when the template is instantiated.
```

```
parameters:   # All the configurable parameters of the template - can include
                # auto-generated passwords etc.
```

The Metadata Section

```
apiVersion: template.openshift.io/v1
kind: Template
```

```
metadata:
```

name: redis-template	# The unique template name
annotations:	# This subsection lists detail of for the # template. There are lots of optional fields # that can be used for template registries.
description: "Description"	# A description with enough detail for users # to understand what is being deployed, links # to additional info etc. Can have newlines
iconClass: "icon-redis"	# An icon for the Web Console (from a list)
tags: "database,nosql"	# Tags for searching and grouping

```
message:
[...]
```


The Message Section

```
apiVersion: template.openshift.io/v1
kind: Template
metadata:

message: "Your password is ${REDIS_PASSWORD}" # Parameters will be filled in at
                                              # instantiation.

labels:
objects:
parameters:
```

The Labels Section

```
apiVersion: template.openshift.io/v1
kind: Template
metadata:
  message:

labels:                                # These labels are applied to each object created when
                                      # the template is instantiated. The labels can be
                                      # parameterized.

  template: redis-template
  redis: master
  app: ${APP}

objects:
parameters:
```

The Parameters Section

```
apiVersion: template.openshift.io/v1
kind: Template
metadata:
message:
labels:
objects:

parameters:                                # a list of configurable parameters
- name: APP                                # a simple parameter with a fixed value
  description: The example app
  value: example
- name: REDIS_PASSWORD                     # an auto-generated 8 character password
  description: Password used for Redis authentication
  from: '[A-Z0-9]{8}'
  generate: expression
```

The Objects Section

```
objects:                                     # A list of objects to be instantiated
- apiVersion: v1                             # In this case only instantiate a single Pod
  kind: Pod
  metadata:
    name: redis-master
  spec:
    containers:
    - env:
      - name: REDIS_PASSWORD
        value: ${REDIS_PASSWORD}             # The environment variable is taken from the
                                              # parameters list.

      image: dockerfile/redis
      name: master
      ports:
      - containerPort: 6379
        protocol: TCP
```

Template and Code Reuse

- The example above is a very simple templates, more sophisticated templates can produce any deployment you can think of.
- Templates can become a useful code reuse tool - they can be checked into Git and versioned.
- For more information see the OpenShift [documentation](#)

Lab 4: Creating and Populating a DB from a Template

OPERATORS

Operators

- An Operator is an application in OpenShift
 - It is an application that manages another application (the **operand** - usually a backend service).
- Operators extend OpenShift by teaching it how to manage your application.
- An Operator makes a service self-managing.
 - Install, upgrade, keep the service running, track metrics etc.

Why Use Operators?

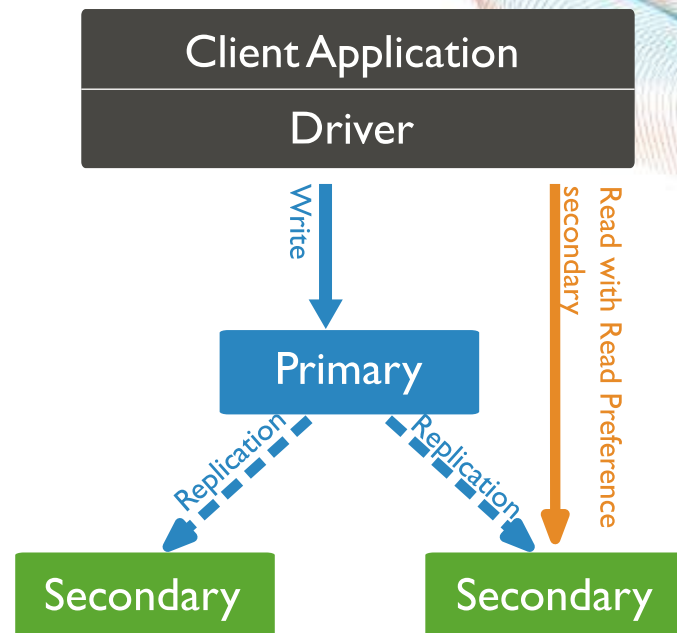
- Lower the barrier to using your service:
 - Skill, time or interest for manual manage are not necessary.
- Services in the main OpenShift service catalogues (OperatorHub.io, RedHat marketplace) *must* have Operators.
- Services in other service catalogues will also benefit from having an Operator.

Stateful apps are a problem in Kubernetes

- K8s manages stateless apps
 - Pod replicas are the same, controllers work for all apps.
- With stateful apps, some pods are different and use unique resources (e.g. Persistent Volumes, PVCs)
- Site-reliability engineers manage stateful apps.

Stateful App Example: MongoDB

- All Pods are *not* the same.
 - One Pod is the Primary, and all others are Secondary Pods federated into the Primary.
 - If the Primary fails, one of the Secondary is promoted and all remaining Secondaries are re-federated into the new Primary.
- Normally a Site Reliability Engineer would be responsible for this.



<https://www.mongodb.com/docs/manual/replication/>

Operators Help with Stateful Apps

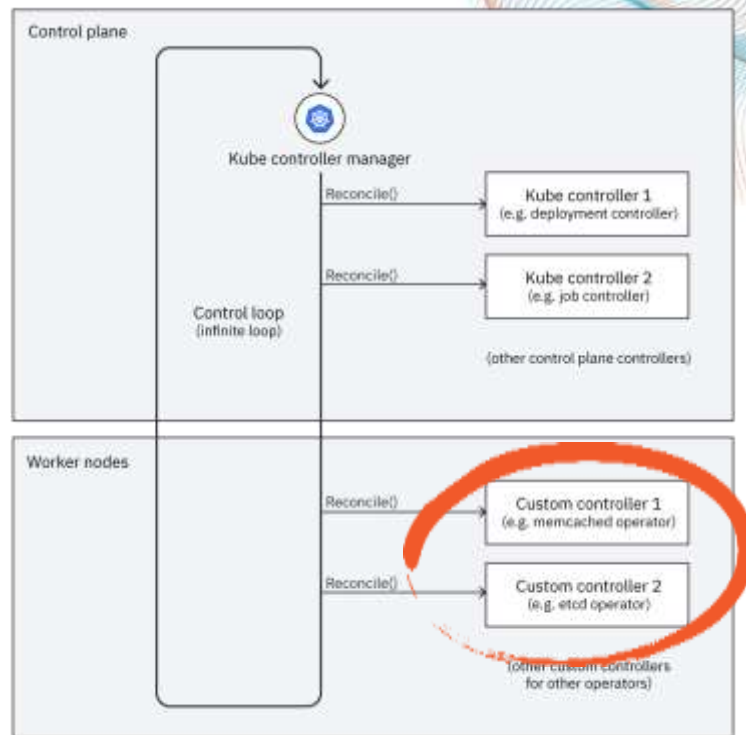
- The Operator knows how the stateful app works and make them more self-managing.
- One Operator can manage multiple operands of the same type.
- Each Operand is configured through variables.
- Operators automate a site reliability engineer's grunt work.

Operator Structure

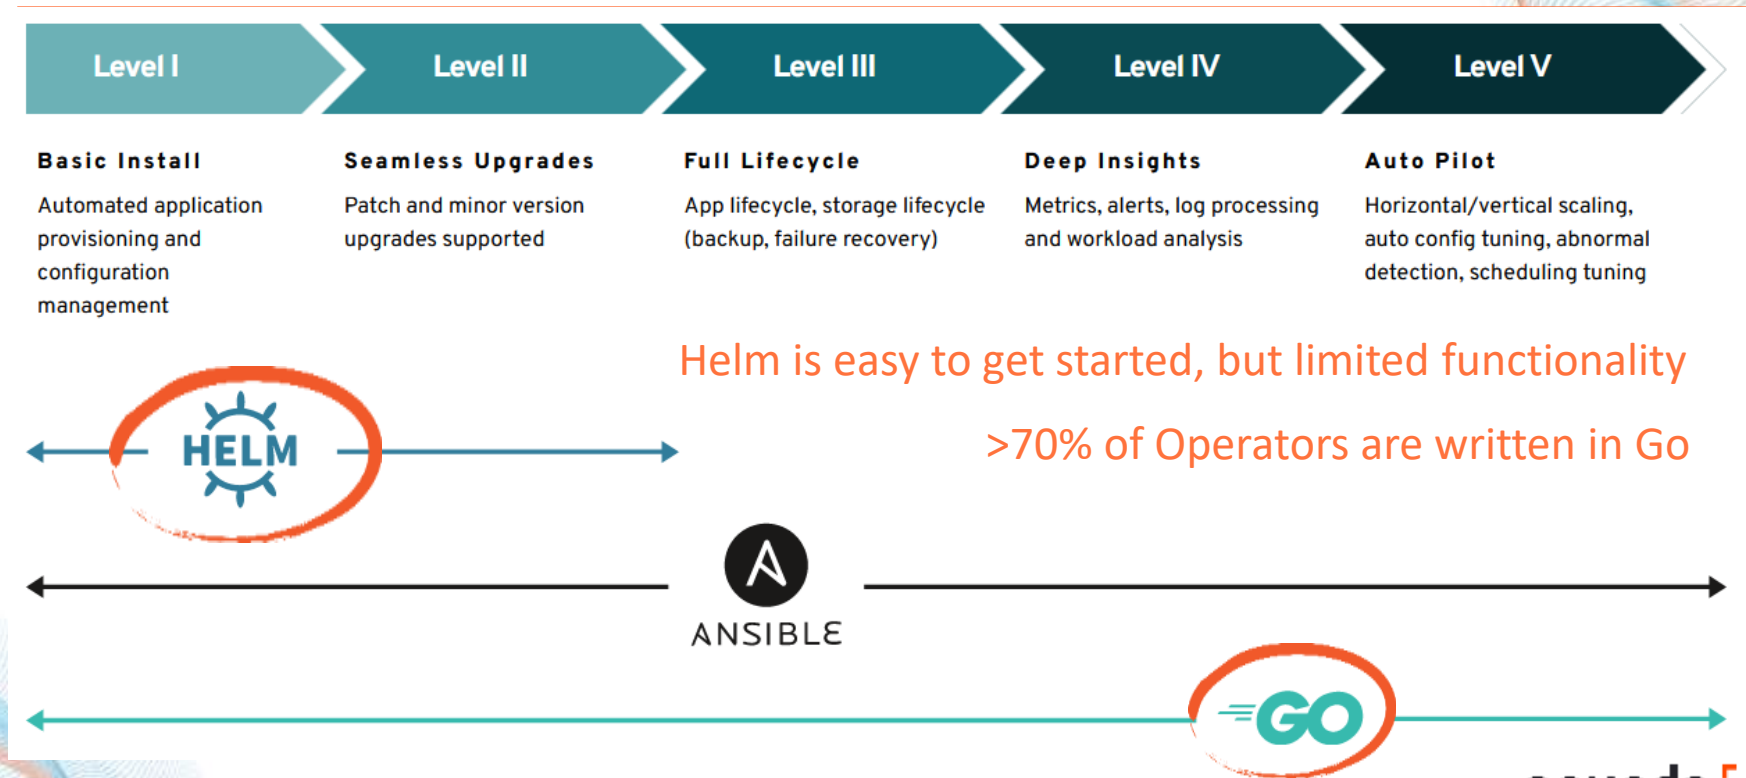
- An Operator runs as an Image in a Container like other applications
- The Operator heart is a Controller, and the main component of the Controller is the `Reconcile()` method.
- The Operator defines its own K8s kind through a Custom Resource Definition (CRD).
- Custom Resources are instances of the CRD

Operators Extend Kubernetes

- In normal running, the Kube controller manager runs an infinite loop calling the `Reconcile()` method of the control plane controllers
 - Checks jobs, deployments etc are matching their declarations.
- Operators just hook onto this loop!



Supported Operator Languages and Control Levels



The Operator Lifecycle Manager

- In OpenShift an Operator Subscription is added for each (desired) Operator from the OperatorHub.
- Operators have their own Operator in OpenShift - the Operator Lifecycle Manager (OLM) - installed by default from OpenShift 4.10.
 - This keeps track of managing the installed Operators.

Operators Conclusion

- Operators shepherd foundation services with custom logic.
- It means adding a database, or a message queue or other common backend services is similar to using a managed cloud service.
- If you are thinking to supply such a service, it is highly recommended you accompany it with an Operator!

Summary

- OpenShift Templates
- Templates in the Web Console and the CLI
- Template File Structure
- Kubernetes Operators
- Operators Structure and operation

Questions and Comments?

